


How close are AI systems to the ceiling of their benchmarks?



— Perfect score (hard ceiling) - - - Label-error ceiling (MMLU) - - - Human baseline - - - Random guessing -◆- Reported result

Scores are as reported by the authors or developers; settings differ (see [data/ai_benchmarks.csv](#)).